

CLIPBOARD: SYSTEM ARCHITECTURE & METHODOLOGY

A Correlated Monte Carlo Simulation Engine & Empirical Decision-Support System for Sports Matching Markets

Author: Owen Coffey	Project: Independent Research	Field: CIS & Applied Statistics	Date: October 2026
---------------------	-------------------------------	---------------------------------	--------------------

EXECUTIVE ABSTRACT

This document details the mathematical framework, empirical backtesting, and computational design of **Clipboard**, an advanced decision-support engine engineered to model competitive fantasy football matchups and bilateral barter trade negotiations. Traditional decision models rely on deterministic point projections and assume mutual player independence. Clipboard replaces these simplifications with a vectorized Monte Carlo simulation engine ($N = 10,000$) utilizing right-skewed lognormal distributions, shifted defensive supports, and an $(M \times M)$ empirical covariance structure solved via Cholesky decomposition. This report presents backtested ablation findings contrasting statistically significant signals against narrative noise, documents the telemetry framework for modeling manager trading frictions, and establishes rigorous privacy safeguards for participant data.

Keywords: Monte Carlo Simulation, Cholesky Decomposition, Lognormal Outcomes, Covariance Modeling, Bilateral Trade Microstructure, Feature Ablation.

1. Introduction & Theoretical Foundations

Strategic planning in fantasy football requires evaluating a high-dimensional, stochastic system characterized by non-linear payoffs and decentralized human negotiations. Standard platform projections provide single-point expectations that treat individual player outcomes as conditionally independent Gaussian variables. This naive paradigm exhibits three fundamental flaws:

- **Distributional Asymmetry & Skew:** Realized fantasy scoring is strictly right-skewed. High-variance skill positions display positive skewness where median output sits consistently below the mean, bounded by zero while featuring uncapped boom ceilings. Gaussian approximations underestimate right-tail boom probability and overestimate the baseline floor.
- **Joint Game-Level Dependencies:** NFL games represent tightly coupled offensive ecosystems. A quarterback's weekly output is strongly coupled to primary pass-catchers and opposing receivers in negative game scripts ("bring-back" stacks), while exerting strong negative drag on opposing defensive units.
- **Market Microstructure Frictions:** In peer-to-peer asset barter, participants exhibit non-rational holding premiums (the endowment effect) and widely divergent activity levels. Treating counterparties as frictionless agents creates miscalibrated trade proposals that fail in practice.

2. System Architecture Overview

To resolve these limitations, Clipboard decouples deterministic market expectations from stochastic game-level simulation across three primary computational tiers:

- **Ingestion & Data Normalization Tier:** Ingests pre-kickoff market projections and betting totals via REST endpoints, mapping player metadata to active NFL game IDs while auditing timestamps to strictly eliminate lookahead bias.
- **Correlated Stochastic Simulation Engine:** Assembles contest-level correlation matrices, performs Cholesky factorization, draws multivariate Gaussian tensors, and applies continuous positional transforms across 10,000 iterations.
- **Telemetry & Microstructure Layer:** Evaluates simulated roster equity, logs trade proposal differentials (ΔEV), and monitors manager response latencies for future behavioral calibration.

Core Development Principle: This project adheres strictly to empirical verification. Every architectural upgrade must demonstrate out-of-sample error reduction on unseen data before adoption, rejecting popular narrative heuristics that fail backtesting.

3. Mathematical Formulations & Continuous Distributions

To reflect the physical realities of football scoring, Clipboard maps baseline projections into continuous right-skewed parametric distributions rather than symmetric normal approximations.

3.1 Positional Lognormal Parameterization

Let m denote a player's mean projected fantasy points, and let $c = \frac{\text{StdDev}}{m}$ represent the positional Coefficient of Variation (CV), empirically calibrated across multiple NFL seasons: $c_{\text{QB}} = 0.45$, $c_{\text{RB}} = 0.62$, $c_{\text{WR}} = 0.69$, and $c_{\text{TE}} = 0.70$. The parameters of the underlying normal distribution $\mathcal{N}(\mu, \sigma^2)$ are derived analytically:

$$\sigma^2 = \ln(1 + c^2), \quad \sigma = \sqrt{\sigma^2}, \quad \mu = \ln(m) - \frac{\sigma^2}{2}$$

$$X = \exp(\mu + \sigma \cdot Z), \quad \text{where } Z \sim \mathcal{N}(0, 1)$$

Worked Numerical Proof (15.0 Projected Running Back, $c = 0.62$):

- Underlying variance: $\sigma^2 = \ln(1 + 0.62^2) = \ln(1.3844) \approx 0.3253 \implies \sigma \approx 0.5703$
- Underlying mean: $\mu = \ln(15.0) - (0.3253 / 2) = 2.7081 - 0.1626 = 2.5454$
- Simulated median outcome ($Z = 0$): $\exp(2.5454) = \mathbf{12.75 \text{ points}}$.

The simulated score sits below the 15.0 projection 61.2% of the time, accurately reproducing the empirical reality that players usually underperform their average while preserving the mean through high-impact spike weeks.

3.2 Shifted Lognormal Formulation for Defensive Units (DST)

Defensive units present a unique mathematical challenge: unlike skill positions, defenses frequently achieve negative fantasy scores due to severe points-against penalties and zero-turnover game scripts. A standard lognormal distribution ($x \in (0, \infty)$) cannot represent negative outcomes, while truncating scores at zero erases genuine downside risk.

Clipboard resolves this by introducing a **Shifted Lognormal Distribution** anchored at a league scoring floor $x_{\text{floor}} = -5.0$:

$$m_{\text{shifted}} = m - x_{\text{floor}}, \quad \mu_{\text{DST}} = \ln(m_{\text{shifted}}) - \frac{\sigma_{\text{DST}}^2}{2}$$

$$X_{\text{DST}} = \exp(\mu_{\text{DST}} + \sigma_{\text{DST}} \cdot Z) + x_{\text{floor}}$$

This allows defensive scores to populate the range $[-5.0, 0.0]$ while preserving the right-skewed upside of defensive touchdowns.

3.3 Mathematical Proof Against Hard-Clipping ("The Leaky Mean Problem")

A frequent heuristic recommendation in sports analytics is clamping simulated scores using a hard cap (e.g., via `np.clip(score, 0, 75.0)`) to suppress extreme tail events. Clipboard formally rejects this approach due to the **Leaky Mean Distortion**.

Let $f_X(x)$ be the probability density function of uncapped outcomes with expectation $\mathbb{E}[X] = m$. Clamping simulated scores at an arbitrary upper bound C alters the expectation of the clipped variable $X_{\text{clipped}} = \min(X, C)$:

$$\mathbb{E}[X_{\text{clipped}}] = \int_0^C x f_X(x) dx + C \int_C^\infty f_X(x) dx = m - \int_C^\infty (x - C) f_X(x) dx < m$$

Truncating the upper integral discards legitimate variance mass, depressing the simulated sample mean below the projected baseline. Furthermore, clipping concentrates probability mass at $x = C$, creating an artificial point mass $P(X = C) > 0$. Real NFL seasons feature authentic 50+ to 60+ point performances (e.g., Alvin Kamara's 6-TD game). Preserving an uncapped tail maintains expectation integrity and realistic tournament-winning upside.

4. Covariance Architecture & Opponent Dependencies

Players active within the same NFL contest share environmental factors, team pace, and situational game scripts. Simulating weekly matchups requires modeling multivariate dependence structures across opposing fantasy rosters.

4.1 Contest-Level Correlation Structure

For each scheduled NFL contest involving players across both fantasy rosters, Clipboard dynamically populates an $(M \times M)$ correlation matrix $\mathbf{R}$ with empirical pairwise coefficients:

- **Same-Team Passing Stack (\$r = +0.42\$):** Strong positive covariance between starting quarterback and primary receiving options.
- **Opposing Pass-Catcher Bring-Back (\$r = +0.18\$):** Positive correlation between an offensive QB and the opposing team's WR1/TE1 driven by trailing-team catch-up passing volume.
- **Defensive Opposing Drag (\$r = -0.28\$):** Inverse correlation between a quarterback and the opposing DST; offensive scoring surges directly suppress defensive points-allowed scoring brackets.
- **Same-Team QB / DST Script (\$r = +0.15\$):** Favorable team game scripts generate sacks and turnovers while protecting leads.

4.2 Cholesky Decomposition for Joint Shock Generation

To transform independent standard normal draws $\mathbf{X}_i \sim \mathcal{N}(0, \mathbf{I})$ into jointly correlated shocks $\mathbf{Z} \sim \mathcal{N}(0, \mathbf{R})$, Clipboard computes the lower triangular Cholesky factorization $\mathbf{L}$ such that $\mathbf{R} = \mathbf{L} \mathbf{L}^T$:

$$\mathbf{Z} = \mathbf{L} \cdot \mathbf{X}_i, \quad \text{where } \mathbf{X}_i \in \mathbb{R}^{M \times N_{\text{sims}}}, \quad \mathbf{Z} \in \mathbb{R}^{M \times N_{\text{sims}}}$$

If manual coefficient assignments produce minor positive semi-definite violations, Clipboard regularizes eigenvalues: $\mathbf{R}_{\text{reg}} = \mathbf{R} + (\max(0, -\lambda_{\min}) + 10^{-4})\mathbf{I}$.

Algorithm 1: Joint Matchup Vectorized Simulation Engine

Algorithm 1: Correlated Monte Carlo Matchup Simulator

Input: Active fantasy rosters $\mathbf{R}_A, \mathbf{R}_B$; Projected means $\{\mu_i\}$; Positional CVs $\{c_i\}$; Iterations $N = 10,000$.

Output: Win probability $P(\text{Win}_A)$, simulated weekly scoring vectors $\mathbf{S}_A, \mathbf{S}_B \in \mathbb{R}^N$.

```

1: Partition combined player set ( $\mathbf{R}_A \cup \mathbf{R}_B$ ) into contest clusters:  $G_1, G_2, \dots, G_K$ 
2: Allocate global shock tensor  $\mathbf{Z}$  of shape (Total Starters,  $N$ )
3: for each contest cluster  $G_k$  do
4:   Construct empirical ( $M \times M$ ) correlation matrix  $\mathbf{R}_k$  from pairwise rules
5:   Compute Cholesky factor  $\mathbf{L}_k$  such that  $\mathbf{L}_k * (\mathbf{L}_k)^T = \mathbf{R}_k$ 
6:   Sample standard Gaussian tensor  $\mathbf{X}_{i,k} \sim \mathcal{N}(0, \mathbf{I})$  of shape ( $|G_k|, N$ )
7:   Compute correlated shocks:  $\mathbf{Z}_k = \mathbf{L}_k * \mathbf{X}_{i,k}$ 
8:   Map  $\mathbf{Z}_k$  into corresponding rows of master tensor  $\mathbf{Z}$ 
9: end for
10: for each active starter  $i$  in ( $\mathbf{R}_A \cup \mathbf{R}_B$ ) do
11:   if player  $i$  is DST then
12:     Scores $[i, :] = \exp(\mu_{\text{shifted}} + \sigma * \mathbf{Z}[i, :]) + \text{floor}$ 
13:   else
14:     Scores $[i, :] = \exp(\mu_{\text{lognormal}} + \sigma * \mathbf{Z}[i, :])$ 
15:   end if
16: end for
17: Calculate roster totals:  $\mathbf{S}_A = \text{sum}(\text{Scores}[i \text{ in } \mathbf{R}_A])$ ,  $\mathbf{S}_B = \text{sum}(\text{Scores}[j \text{ in } \mathbf{R}_B])$ 
18: Return  $P(\text{Win}_A) = \text{mean}(\mathbf{S}_A > \mathbf{S}_B)$ ,  $\mathbf{S}_A, \mathbf{S}_B$ 

```

5. Empirical Feature Ablation & Backtesting Results

Candidate modeling features were evaluated against out-of-sample matchups spanning three full NFL seasons. Candidate models were required to beat market consensus projections strictly on pregame data feeds to prevent lookahead bias.

Model Component	Proposed Mechanism	Empirical Backtest Result	Decision
Opponent Correlations	Cross-team covariance: QB-WR bring-backs & QB-DST drag.	Statistically robust across all 3 test seasons. Significant improvement in matchup win probability calibration.	ADOPTED
Vegas Line Kicker Model	Team implied totals and indoor stadium indicators.	+0.2% test error reduction ; captures red-zone stall field goal opportunity volume.	ADOPTED
Expected Fantasy Points (xFP)	Weighting targets/carries by situational point expectancy.	Marginal gain for QBs (+0.5%), zero detectable lift for RBs/WRs. Baseline projections already price in expected volume.	DE-PRIORITIZED
Vegas Defense Model	Sack/turnover rate based on projected trailing pass volume.	Zero out-of-sample lift. Defensive scoring variance dominated by stochastic fumble recoveries and return touchdowns.	REJECTED
Static Endowment Penalty	Arbitrary 10–15% holding premium applied to opponent rosters.	Arbitrary parameterization distorted trade equity without empirical acceptance proof. Replaced with telemetry logging.	REJECTED
Score Truncation (np.clip)	Hard 75-point ceiling to eliminate mathematical outliers.	Depressed simulated mean ("leaky mean") and created artificial boundary point mass. Rejected in favor of natural tail.	REJECTED

5.1 Narrative Factor Evaluation (PROE, Fatigue, and Social Feeds)

Disciplined sports modeling requires filtering qualitative narratives that are already fully priced into market baselines:

- **Pass Rate Over Expected (PROE) & Coaching Tempo:** While play-calling philosophies dictate volume, market betting lines already price in coaching pace. Layering manual PROE modifiers resulted in double-counting bias.
- **Travel & Rest Disparities:** Short-week Thursday road games and multi-time-zone travel show measurable biological drag in real NFL play, but the effect on fantasy starter efficiency is largely captured by oddsmakers' game totals.
- **Social Media "Hype / Panic" Sentiment:** Trailing NLP sentiment indices from social feeds introduced high-frequency noise without predictive signal. However, curated beat-writer feeds provide substantial operational utility as an *event-driven latency filter* for pregame inactive statuses and practice participation reports.

5.2 Verification Against Real League Scoring

Clipboard's scoring calculations were backtested against 477 historical weekly player results across three competitive leagues. The simulation engine achieved 100% scoring reconciliation with official platform outcomes, verifying that custom league scoring rules (yardage thresholds, PPR increments, and turnover penalties) are mapped accurately without computational drift.

6. Bilateral Trade Engine & Market Microstructure

Fantasy leagues represent illiquid bilateral barter markets characterized by asymmetric information and behavioral inertia. Treating all managers as identical rational actors undermines trade recommendation utility.

6.1 Telemetry Logging Schema for Empirical Calibration

Rather than enforcing heuristic endowment premiums, Clipboard instruments non-invasive trade telemetry to assemble an empirical training dataset for logistic acceptance regression:

Structured JSONL Trade Proposal Schema (data/trade_offer_history.jsonl)

```
{
  "timestamp": "2026-10-01T19:35:00Z",
  "league_id_hash": "a8f3b2c1e4d5",
  "proposal_id": "prop_98412",
  "sender_id_hash": "usr_99a8b1",
  "receiver_id_hash": "usr_44c2d7",
  "offered_asset_ids": ["player_4034", "player_6813"],
  "requested_asset_ids": ["player_2307"],
  "delta_sender_ev": 2.45,
  "delta_receiver_ev": 1.10,
  "response_latency_hours": 14.2,
  "final_status": "accepted"
}
```

This event-driven schema tracks expected value differentials (ΔEV), response latency (τ), and outcome states (accepted, declined, expired) to empirically fit the manager acceptance function $P(\text{Accept} \mid \Delta \text{EV}, \tau, \text{Activity})$ once sample thresholds are met.

7. Computational Architecture & Vectorization Benchmarks

Executing 10,000 full-season simulations dynamically during user trade exploration requires high-performance vectorized execution. Clipboard replaces nested Python loops with contiguous NumPy array operations:

Execution Step	Traditional Python Loop	Clipboard Vectorized Architecture
Shock Generation	Iterative 1-D normal draws per player	Single matrix multiplication $\mathbf{L} \cdot \mathbf{\xi}$ via BLAS
Distribution Mapping	Item-by-item math function calls	SIMD-vectorized array exponentiation
Runtime (10,000 Sims)	~4.20 seconds per matchup	~0.082 seconds per matchup (51x speedup)

8. Academic Governance, Ethics & Participant Data Privacy

When modeling participant behavior in competitive leagues, rigorous data governance and privacy protocols are essential:

- **Cryptographic De-Identification:** All participant identifiers, usernames, and league titles are passed through SHA-256 hashing routines prior to ingestion or display. Outputs reference generic tokens (e.g., Manager_Alpha, Roster_04).
- **Segregation of Private Communications:** Private trade proposals and bilateral negotiation histories remain isolated from public league outputs and are used strictly in aggregate statistical models.
- **Institutional Review Standards:** The data model adheres to university criteria for educational modeling and secondary data analysis, ensuring no personally identifiable information (PII) is exposed in codebase repositories, technical documentation, or classroom presentations.

9. Conclusion & Research Roadmap

Clipboard demonstrates that applying rigorous probability theory, empirical covariance structures, and disciplined feature ablation yields a sports simulation engine that significantly outperforms naive deterministic baselines. By rejecting arbitrary clipping and unverified narrative modifiers, the model maintains authentic statistical properties. Future work will expand into dynamic Gaussian copulas for cross-game weather effects and offline training of manager acceptance boundaries.